

Deploy a registry server

Estimated reading time: 18 minutes

Before you can deploy a registry, you need to install Docker on the host. A registry is an instance of the `registry` image, and runs within Docker.

This topic provides basic information about deploying and configuring a registry. For an exhaustive list of configuration options, see the configuration reference (<https://docs.docker.com/registry/configuration/>).

If you have an air-gapped datacenter, see Considerations for air-gapped registries (</registry/deploying/#considerations-for-air-gapped-registries>).

Run a local registry

Use a command like the following to start the registry container:

```
$ docker run -d -p 5000:5000 --restart=always --name registry registry:2
```

The registry is now ready to use.

Warning: These first few examples show registry configurations that are only appropriate for testing. A production-ready registry must be protected by TLS and should ideally use an access-control mechanism. Keep reading and then continue to the configuration guide (<https://docs.docker.com/registry/configuration/>) to deploy a production-ready registry.

Copy an image from Docker Hub to your registry

You can pull an image from Docker Hub and push it to your registry. The following example pulls the `ubuntu:16.04` image from Docker Hub and re-tags it as `my-ubuntu`, then pushes it to the local registry. Finally, the `ubuntu:16.04` and `my-ubuntu` images are deleted locally and the `my-ubuntu` image is pulled from the local registry.

1. Pull the `ubuntu:16.04` image from Docker Hub.

```
$ docker pull ubuntu:16.04
```

